

distillbert-for-sentiment-analysis

November 21, 2024

```
[ ]: import pandas as pd
import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader, random_split
from transformers import DistilBertModel, DistilBertTokenizerFast
from torch.utils.data import DataLoader
from transformers import AdamW
```

Preparing dataset

```
[ ]: !wget https://raw.githubusercontent.com/kyuz0/llm-chronicles/main/datasets/
↪restaurant_reviews.csv
```

```
--2024-11-13 09:35:33-- https://raw.githubusercontent.com/kyuz0/llm-
chronicles/main/datasets/restaurant_reviews.csv
Resolving raw.githubusercontent.com (raw.githubusercontent.com)...
185.199.108.133, 185.199.109.133, 185.199.110.133, ...
Connecting to raw.githubusercontent.com
(raw.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 2861025 (2.7M) [text/plain]
Saving to: 'restaurant_reviews.csv'

restaurant_reviews. 100%[=====>] 2.73M --.-KB/s in 0.05s

2024-11-13 09:35:34 (54.0 MB/s) - 'restaurant_reviews.csv' saved
[2861025/2861025]
```

```
[ ]: # Load the dataset
df = pd.read_csv('restaurant_reviews.csv')

# Map sentiments to numerical labels
sentiment_mapping = {'negative': 0, 'neutral': 1, 'positive': 2}
df['Rating'] = df['Rating'].map(sentiment_mapping)

[ ]: df.head()
```

```
[ ]:
```

	Review	Rating
0	The ambience was good food was quite good . ha...	2
1	Ambience is too good for a pleasant evening. S...	2
2	A must try.. great food great ambience. Thnx f...	2
3	Soumen das and Arun was a great guy. Only beca...	2
4	Food is good.we ordered Kodi drumsticks and ba...	2

```
[ ]: # Display the first few rows of the dataframe
print(df.head())

# Display statistics about the dataset
print("\nDataset Statistics:")
print(df['Rating'].value_counts())
```

	Review	Rating
0	The ambience was good food was quite good . ha...	2
1	Ambience is too good for a pleasant evening. S...	2
2	A must try.. great food great ambience. Thnx f...	2
3	Soumen das and Arun was a great guy. Only beca...	2
4	Food is good.we ordered Kodi drumsticks and ba...	2

Dataset Statistics:

Rating

2 6331

0 2428

1 1192

Name: count, dtype: int64

PyTorch Dataset and Dataloader

```
[ ]: class ReviewDataset(Dataset):
    def __init__(self, csv_file, tokenizer, max_length):
        self.dataset = pd.read_csv(csv_file)
        self.tokenizer = tokenizer
        self.max_length = max_length
        # Map sentiments to numerical labels
        self.label_dict = {'negative': 0, 'neutral': 1, 'positive': 2}

    def __len__(self):
        return len(self.dataset)

    def __getitem__(self, idx):
        review_text = self.dataset.iloc[idx, 0] # Assuming reviewText is the
        ↪first column
        sentiment = self.dataset.iloc[idx, 1] # Assuming sentiment is the
        ↪second column
        labels = self.label_dict[sentiment] # Convert sentiment to numerical
        ↪label
```

```

        # Tokenize the review text
        encoding = self.tokenizer.encode_plus(
            review_text,
            add_special_tokens=True, # Add [CLS] token at the start for
↪classification
            max_length=self.max_length,
            return_token_type_ids=False,
            padding='max_length',
            return_attention_mask=True,
            return_tensors='pt',
            truncation=True
        )

        return {
            'review_text': review_text,
            'input_ids': encoding['input_ids'].flatten(),
            'attention_mask': encoding['attention_mask'].flatten(), # this is NOT
↪self-attention!
            'labels': torch.tensor(labels, dtype=torch.long)
        }

```

```

[ ]: tokenizer = DistilBertTokenizerFast.from_pretrained('distilbert-base-uncased')
review_dataset = ReviewDataset('restaurant_reviews.csv', tokenizer, 512)

```

```

/usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/_auth.py:94:
UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab
(https://huggingface.co/settings/tokens), set it as secret in your Google Colab
and restart your session.
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access
public models or datasets.

```

```

warnings.warn(

tokenizer_config.json:  0%|          | 0.00/48.0 [00:00<?, ?B/s]
vocab.txt:  0%|          | 0.00/232k [00:00<?, ?B/s]
tokenizer.json:  0%|          | 0.00/466k [00:00<?, ?B/s]
config.json:  0%|          | 0.00/483 [00:00<?, ?B/s]

```

```

[ ]: review_dataset[0]

```

```

[ ]: {'review_text': 'The ambience was good food was quite good . had Saturday lunch
which was cost effective . Good place for a sate brunch. One can also chill with
friends and or parents. Waiter Soumen Das was really courteous and helpful.',

```


[illegible]

```
[ ]: from torch.utils.data import DataLoader, random_split

# Split dataset into training and validation
train_size = int(0.8 * len(df))
val_size = len(df) - train_size
train_dataset, test_dataset = random_split(review_dataset, [train_size,
    ↪val_size])

train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=16)
```

```
[ ]: # Show number of batches
len(train_loader), len(test_loader)
```

[]: (498, 125)

Fine-tuning with custom classifier layer

```
[ ]: class CustomDistilBertForSequenceClassification(nn.Module):
    def __init__(self, num_labels=3):
        super(CustomDistilBertForSequenceClassification, self).__init__()
        self.distilbert = DistilBertModel.
        ↪from_pretrained('distilbert-base-uncased') #base model from hugging face
        self.pre_classifier = nn.Linear(768, 768) # DistilBERT's hidden size,
        ↪is 768
        self.dropout = nn.Dropout(0.3)
        self.classifier = nn.Linear(768, num_labels)

    def forward(self, input_ids, attention_mask):
        distilbert_output = self.distilbert(input_ids=input_ids,
        ↪attention_mask=attention_mask)
        hidden_state = distilbert_output[0] # (batch_size, sequence_length,
        ↪hidden_size)
        pooled_output = hidden_state[:, 0] # we take the representation of the
        ↪[CLS] token (first token)
        pooled_output = self.pre_classifier(pooled_output)
        pooled_output = nn.ReLU()(pooled_output)
        pooled_output = self.dropout(pooled_output) # regularization
        logits = self.classifier(pooled_output)
        return logits
```

```
[ ]: model = CustomDistilBertForSequenceClassification()
```

```
model.safetensors: 0%|          | 0.00/268M [00:00<?, ?B/s]
```

```
[ ]: # Inspect DistilBERT
print(model.distilbert)
```

```
DistilBertModel(
  (embeddings): Embeddings(
    (word_embeddings): Embedding(30522, 768, padding_idx=0)
    (position_embeddings): Embedding(512, 768)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
  )
  (transformer): Transformer(
    (layer): ModuleList(
      (0-5): 6 x TransformerBlock(
        (attention): DistilBertSdpaAttention(
          (dropout): Dropout(p=0.1, inplace=False)
          (q_lin): Linear(in_features=768, out_features=768, bias=True)
          (k_lin): Linear(in_features=768, out_features=768, bias=True)
          (v_lin): Linear(in_features=768, out_features=768, bias=True)
          (out_lin): Linear(in_features=768, out_features=768, bias=True)
        )
      )
    )
  )
)
```

```

(sa_layer_norm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
(ffn): FFN(
  (dropout): Dropout(p=0.1, inplace=False)
  (lin1): Linear(in_features=768, out_features=3072, bias=True)
  (lin2): Linear(in_features=3072, out_features=768, bias=True)
  (activation): GELUActivation()
)
(output_layer_norm): LayerNorm((768,), eps=1e-12,
elementwise_affine=True)
)
)
)
)
)

```

Fine-tuning

```

[ ]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model.to(device)

optimizer = AdamW(model.parameters(), lr=5e-5)

model.train()
for epoch in range(5):
    for i, batch in enumerate(train_loader):
        input_ids = batch['input_ids'].to(device)
        attention_mask = batch['attention_mask'].to(device)
        labels = batch['labels'].to(device)

        optimizer.zero_grad()
        logits = model(input_ids=input_ids, attention_mask=attention_mask)
        loss = nn.CrossEntropyLoss()(logits, labels)
        loss.backward()
        optimizer.step()

        if (i + 1) % 100 == 0:
            print(f"Epoch {epoch + 1}, Batch {i + 1}, Loss: {loss.item():.4f}")

```

```

Epoch 1, Batch 100, Loss: 0.4197
Epoch 1, Batch 200, Loss: 0.3587
Epoch 1, Batch 300, Loss: 0.1825
Epoch 1, Batch 400, Loss: 0.1685
Epoch 2, Batch 100, Loss: 0.2091
Epoch 2, Batch 200, Loss: 0.2948
Epoch 2, Batch 300, Loss: 0.1590
Epoch 2, Batch 400, Loss: 0.0778
Epoch 3, Batch 100, Loss: 0.1450
Epoch 3, Batch 200, Loss: 0.1897
Epoch 3, Batch 300, Loss: 0.2143

```



```
Epoch 3, Batch 400, Loss: 0.0392
Epoch 4, Batch 100, Loss: 0.4496
Epoch 4, Batch 200, Loss: 0.0672
Epoch 4, Batch 300, Loss: 0.1173
Epoch 4, Batch 400, Loss: 0.0902
Epoch 5, Batch 100, Loss: 0.0152
Epoch 5, Batch 200, Loss: 0.0489
Epoch 5, Batch 300, Loss: 0.0172
Epoch 5, Batch 400, Loss: 0.0088
```

**** Evaluation****

```
[ ]: model.eval()
total_correct = 0
total = 0
for batch in test_loader:
    input_ids = batch['input_ids'].to(device)
    attention_mask = batch['attention_mask'].to(device)
    labels = batch['labels'].to(device)

    with torch.inference_mode():
        logits = model(input_ids=input_ids, attention_mask=attention_mask)
        predictions = torch.argmax(logits, dim=1)
        total_correct += (predictions == labels).sum().item()
        total += predictions.size(0)

print(f'Test Accuracy: {total_correct / total:.4f}')
```

Test Accuracy: 0.8553

```
[ ]: def predict_sentiment(review_text, model, tokenizer, max_length = 512):
    """
    Predicts the sentiment of a given review text.

    Args:
    - review_text (str): The review text to analyze.
    - model (torch.nn.Module): The fine-tuned sentiment analysis model.
    - tokenizer (PreTrainedTokenizer): The tokenizer for encoding the text.
    - max_length (int): The maximum sequence length for the model.

    Returns:
    - sentiment (str): The predicted sentiment label ('negative', 'neutral', 'positive').
    """

    # Ensure the model is in evaluation mode
    model.eval()
```

```

# Tokenize the input text
encoding = tokenizer.encode_plus(
    review_text,
    add_special_tokens=True,
    max_length=max_length,
    return_token_type_ids=False,
    padding='max_length',
    return_attention_mask=True,
    return_tensors='pt',
    truncation=True
)

input_ids = encoding['input_ids']
attention_mask = encoding['attention_mask']

# Move tensors to the same device as the model
input_ids = input_ids.to(device)
attention_mask = attention_mask.to(device)

with torch.inference_mode():
    # Forward pass, get logits
    logits = model(input_ids=input_ids, attention_mask=attention_mask)

# Extract the highest scoring output
prediction = torch.argmax(logits, dim=1).item()

# Map prediction to label
label_dict = {0: 'negative', 1: 'neutral', 2: 'positive'}
sentiment = label_dict[prediction]

return sentiment

```

```

[ ]: # Test
review_1 = "We ordered from Papa Johns a so-called pizza... what to say? I'd
    ↪rather eat a piece of dry cardboard, calling this pizza is an insult to
    ↪Italians! "
review_2 = "I guess PizzaHut is decent but far from the Italian pizza. This is
    ↪not going to blow you away, but still quite ok in the end."
review_3 = "Gino's pizza is what authentical Neapolian pizza tastes like,
    ↪highly recommended."

print(predict_sentiment(review_1, model, tokenizer))
print(predict_sentiment(review_2, model, tokenizer))
print(predict_sentiment(review_3, model, tokenizer))

```

negative

neutral

positive

[]: